

Rule v2 decisions per dataset

D1 (R9.4, 0.0 Mb, GC=38%) -> default/static

D2 (R9.4, 5.3 Mb, GC=50%) -> hmm/static

D3 (R9.4, 12.0 Mb, GC=38%) -> hmm/static

D4 (R9.4, 112.0 Mb, GC=64%) -> hmm/static

D6 (R10.4, 5.3 Mb, GC=50%) -> binseg/auto BOOST=1.0

D7 (R10.4, 2.8 Mb, GC=33%) -> default/static


RB_dmel (R10.4.1, 143.0 Mb, GC=42%) -> window/auto BOOST=1.0

RB_ecoli (R10.4.1, 5.3 Mb, GC=50%) -> hmm/static


RB_hsap_sub (R10.4.1, 46.0 Mb, GC=41%) -> window/auto BOOST=1.0

default
binseg
hmm
window

Accuracy: default vs rule v2 vs per-dataset oracle




Wall time per dataset (median; whiskers = min/max across runs)




Per-dataset wall-time speed-up of rule v2 vs default


geo-mean = 0.94x




Accuracy delta: rule v2 - default (positive = better)




Auto-K BOOST: accuracy effect (top-K family, R10.4)



Auto-K BOOST: wall-time effect



Run-to-run reproducibility (top 25 by spread)



All-segmenter Pareto over 8 datasets
o = static, ^ = Auto-K (BOOST=1.0), D = Auto-K + BOOST

